

Data Analytics for house pricing dataset

October 1, 2024

0.1 The Dataset

Variable	Description
id	A notation for a house
date	Date house was sold
price	Price is prediction target
bedrooms	Number of bedrooms
bathrooms	Number of bathrooms
sqft_living	Square footage of the home
sqft_lot	Square footage of the lot
floors	Total floors (levels) in house
waterfront	House which has a view to a waterfront
view	Has been viewed
condition	How good the condition is overall
grade	overall grade given to the housing unit, based on King County grading system
sqft_above	Square footage of house apart from basement
sqft_basement	Square footage of the basement
yr_built	Built Year
yr_renovated	Year when house was renovated
zipcode	Zip code
lat	Latitude coordinate
long	Longitude coordinate
sqft_living15	Living room area in 2015(implies– some renovations) This might or might not have affected the lotsize area
sqft_lot15	LotSize area in 2015(implies– some renovations)

0.2 Import the required libraries

```
[69]: # Surpress warnings:
def warn(*args, **kwargs):
    pass
import warnings
warnings.warn = warn

import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
```

```
import seaborn as sns
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler, PolynomialFeatures
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
%matplotlib inline
```

1 Module 1: Importing Data Sets

```
[25]: filepath='https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/
↳IBMDeveloperSkillsNetwork-DA0101EN-SkillsNetwork/labs/FinalModule_Coursera/
↳data/kc_house_data_NaN.csv'
df = pd.read_csv(filepath)
```

```
[26]: df.head()
```

```
[26]:   Unnamed: 0    id          date    price  bedrooms  bathrooms \
0          0  7129300520  20141013T000000  221900.0         3.0         1.00
1          1  6414100192  20141209T000000  538000.0         3.0         2.25
2          2  5631500400  20150225T000000  180000.0         2.0         1.00
3          3  2487200875  20141209T000000  604000.0         4.0         3.00
4          4  1954400510  20150218T000000  510000.0         3.0         2.00
```

```
   sqft_living  sqft_lot  floors  waterfront  ...  grade  sqft_above \
0          1180      5650      1.0           0  ...      7          1180
1          2570      7242      2.0           0  ...      7          2170
2           770     10000      1.0           0  ...      6           770
3          1960      5000      1.0           0  ...      7          1050
4          1680      8080      1.0           0  ...      8          1680
```

```
   sqft_basement  yr_built  yr_renovated  zipcode    lat    long \
0              0     1955              0    98178  47.5112 -122.257
1             400     1951             1991    98125  47.7210 -122.319
2              0     1933              0    98028  47.7379 -122.233
3             910     1965              0    98136  47.5208 -122.393
4              0     1987              0    98074  47.6168 -122.045
```

```
   sqft_living15  sqft_lot15
0           1340          5650
1           1690          7639
2           2720          8062
3           1360          5000
4           1800          7503
```

```
[5 rows x 22 columns]
```

1.0.1 Question 1

Display the data types of each column using the function `dtypes`. Take a screenshot of your code and output. You will need to submit the screenshot for the final project.

```
[27]: df.dtypes
```

```
[27]: Unnamed: 0      int64
      id           int64
      date         object
      price        float64
      bedrooms     float64
      bathrooms    float64
      sqft_living   int64
      sqft_lot      int64
      floors       float64
      waterfront    int64
      view         int64
      condition    int64
      grade        int64
      sqft_above    int64
      sqft_basement int64
      yr_built      int64
      yr_renovated  int64
      zipcode       int64
      lat          float64
      long         float64
      sqft_living15 int64
      sqft_lot15    int64
      dtype: object
```

We use the method `describe` to obtain a statistical summary of the dataframe.

```
[28]: df.describe()
```

```
[28]:
```

	Unnamed: 0	id	price	bedrooms	bathrooms	\
count	21613.00000	2.161300e+04	2.161300e+04	21600.000000	21603.000000	
mean	10806.00000	4.580302e+09	5.400881e+05	3.372870	2.115736	
std	6239.28002	2.876566e+09	3.671272e+05	0.926657	0.768996	
min	0.00000	1.000102e+06	7.500000e+04	1.000000	0.500000	
25%	5403.00000	2.123049e+09	3.219500e+05	3.000000	1.750000	
50%	10806.00000	3.904930e+09	4.500000e+05	3.000000	2.250000	
75%	16209.00000	7.308900e+09	6.450000e+05	4.000000	2.500000	
max	21612.00000	9.900000e+09	7.700000e+06	33.000000	8.000000	

	sqft_living	sqft_lot	floors	waterfront	view	\
count	21613.000000	2.161300e+04	21613.000000	21613.000000	21613.000000	
mean	2079.899736	1.510697e+04	1.494309	0.007542	0.234303	
std	918.440897	4.142051e+04	0.539989	0.086517	0.766318	

min	290.000000	5.200000e+02	1.000000	0.000000	0.000000
25%	1427.000000	5.040000e+03	1.000000	0.000000	0.000000
50%	1910.000000	7.618000e+03	1.500000	0.000000	0.000000
75%	2550.000000	1.068800e+04	2.000000	0.000000	0.000000
max	13540.000000	1.651359e+06	3.500000	1.000000	4.000000

	...	grade	sqft_above	sqft_basement	yr_built	\
count	...	21613.000000	21613.000000	21613.000000	21613.000000	
mean	...	7.656873	1788.390691	291.509045	1971.005136	
std	...	1.175459	828.090978	442.575043	29.373411	
min	...	1.000000	290.000000	0.000000	1900.000000	
25%	...	7.000000	1190.000000	0.000000	1951.000000	
50%	...	7.000000	1560.000000	0.000000	1975.000000	
75%	...	8.000000	2210.000000	560.000000	1997.000000	
max	...	13.000000	9410.000000	4820.000000	2015.000000	

	yr_renovated	zipcode	lat	long	sqft_living15	\
count	21613.000000	21613.000000	21613.000000	21613.000000	21613.000000	
mean	84.402258	98077.939805	47.560053	-122.213896	1986.552492	
std	401.679240	53.505026	0.138564	0.140828	685.391304	
min	0.000000	98001.000000	47.155900	-122.519000	399.000000	
25%	0.000000	98033.000000	47.471000	-122.328000	1490.000000	
50%	0.000000	98065.000000	47.571800	-122.230000	1840.000000	
75%	0.000000	98118.000000	47.678000	-122.125000	2360.000000	
max	2015.000000	98199.000000	47.777600	-121.315000	6210.000000	

	sqft_lot15
count	21613.000000
mean	12768.455652
std	27304.179631
min	651.000000
25%	5100.000000
50%	7620.000000
75%	10083.000000
max	871200.000000

[8 rows x 21 columns]

2 Module 2: Data Wrangling

2.0.1 Question 2

Drop the columns “id” and “Unnamed: 0” from axis 1 using the method `drop()`, then use the method `describe()` to obtain a statistical summary of the data. Make sure the `inplace` parameter is set to `True`. Take a screenshot of your code and output. You will need to submit the screenshot for the final project.

```
[29]: df.drop(["id", "Unnamed: 0"], axis=1, inplace=True)
```

```
[30]: df.describe()
```

```
[30]:
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot \
count	2.161300e+04	21600.000000	21603.000000	21613.000000	2.161300e+04
mean	5.400881e+05	3.372870	2.115736	2079.899736	1.510697e+04
std	3.671272e+05	0.926657	0.768996	918.440897	4.142051e+04
min	7.500000e+04	1.000000	0.500000	290.000000	5.200000e+02
25%	3.219500e+05	3.000000	1.750000	1427.000000	5.040000e+03
50%	4.500000e+05	3.000000	2.250000	1910.000000	7.618000e+03
75%	6.450000e+05	4.000000	2.500000	2550.000000	1.068800e+04
max	7.700000e+06	33.000000	8.000000	13540.000000	1.651359e+06

	floors	waterfront	view	condition	grade \
count	21613.000000	21613.000000	21613.000000	21613.000000	21613.000000
mean	1.494309	0.007542	0.234303	3.409430	7.656873
std	0.539989	0.086517	0.766318	0.650743	1.175459
min	1.000000	0.000000	0.000000	1.000000	1.000000
25%	1.000000	0.000000	0.000000	3.000000	7.000000
50%	1.500000	0.000000	0.000000	3.000000	7.000000
75%	2.000000	0.000000	0.000000	4.000000	8.000000
max	3.500000	1.000000	4.000000	5.000000	13.000000

	sqft_above	sqft_basement	yr_built	yr_renovated	zipcode \
count	21613.000000	21613.000000	21613.000000	21613.000000	21613.000000
mean	1788.390691	291.509045	1971.005136	84.402258	98077.939805
std	828.090978	442.575043	29.373411	401.679240	53.505026
min	290.000000	0.000000	1900.000000	0.000000	98001.000000
25%	1190.000000	0.000000	1951.000000	0.000000	98033.000000
50%	1560.000000	0.000000	1975.000000	0.000000	98065.000000
75%	2210.000000	560.000000	1997.000000	0.000000	98118.000000
max	9410.000000	4820.000000	2015.000000	2015.000000	98199.000000

	lat	long	sqft_living15	sqft_lot15
count	21613.000000	21613.000000	21613.000000	21613.000000
mean	47.560053	-122.213896	1986.552492	12768.455652
std	0.138564	0.140828	685.391304	27304.179631
min	47.155900	-122.519000	399.000000	651.000000
25%	47.471000	-122.328000	1490.000000	5100.000000
50%	47.571800	-122.230000	1840.000000	7620.000000
75%	47.678000	-122.125000	2360.000000	10083.000000
max	47.777600	-121.315000	6210.000000	871200.000000

```
[31]: df.isna().sum()
```

```
[31]: date          0
      price         0
      bedrooms     13
      bathrooms    10
      sqft_living   0
      sqft_lot      0
      floors        0
      waterfront    0
      view          0
      condition     0
      grade         0
      sqft_above    0
      sqft_basement 0
      yr_built      0
      yr_renovated  0
      zipcode       0
      lat           0
      long          0
      sqft_living15 0
      sqft_lot15    0
      dtype: int64
```

We can see we have missing values for the columns bedrooms and bathrooms

```
[32]: print("number of NaN values for the column bedrooms :", df['bedrooms'].isnull().
      ↪sum())
      print("number of NaN values for the column bathrooms :", df['bathrooms'].
      ↪isnull().sum())
```

```
number of NaN values for the column bedrooms : 13
number of NaN values for the column bathrooms : 10
```

```
[33]: mean = df['bedrooms'].mean()
      df.replace(np.nan, mean, inplace=True)
```

```
[35]: mean2 = df['bathrooms'].mean()
      df.replace(np.nan, mean, inplace=True)
```

```
[36]: print("number of NaN values for the column bedrooms :", df['bedrooms'].isnull().
      ↪sum())
      print("number of NaN values for the column bathrooms :", df['bathrooms'].
      ↪isnull().sum())
```

```
number of NaN values for the column bedrooms : 0
number of NaN values for the column bathrooms : 0
```

3 Module 3: Exploratory Data Analysis

3.0.1 Question 3

Use the method `value_counts` to count the number of houses with unique floor values, use the method `.to_frame()` to convert it to a data frame. Take a screenshot of your code and output. You will need to submit the screenshot for the final project.

```
[43]: df['floors'].value_counts().to_frame()
```

```
[43]:
```


	count
1.0	10680
2.0	8241
1.5	1910
3.0	613
2.5	161
3.5	8

3.0.2 Question 4

Use the function `boxplot` in the `seaborn` library to determine whether houses with a waterfront view or without a waterfront view have more price outliers. Take a screenshot of your code and boxplot. You will need to submit the screenshot for the final project.

```
[44]: sns.boxplot ( x= 'waterfront', y= 'price', data=df)
```

```
[44]: <Axes: xlabel='waterfront', ylabel='price'>
```




3.0.3 Question 5

Use the function `regplot` in the `seaborn` library to determine if the feature `sqft_above` is negatively or positively correlated with `price`. Take a screenshot of your code and scatterplot. You will need to submit the screenshot for the final project.

```
[45]: sns.regplot( x='sqft_above', y='price', data=df)
```

```
[45]: <Axes: xlabel='sqft_above', ylabel='price'>
```

```
[49]: df.dtypes
```

```
[49]: date           object
      price         float64
      bedrooms      float64
      bathrooms     float64
      sqft_living    int64
      sqft_lot       int64
      floors         float64
      waterfront     int64
      view           int64
      condition      int64
      grade          int64
      sqft_above     int64
      sqft_basement  int64
      yr_built       int64
      yr_renovated   int64
      zipcode        int64
      lat            float64
      long           float64
```

```
sqft_living15      int64
sqft_lot15         int64
dtype: object
```

```
[50]: # Drop the 'date' column
df_numeric = df.drop(columns=['date'])

# Calculate the correlation matrix for numeric columns only
correlation_matrix = df_numeric.corr()

# Get the correlation values with respect to 'price' and sort them
price_correlation = correlation_matrix['price'].sort_values()
print(price_correlation)
```

```
zipcode      -0.053203
long          0.021626
condition     0.036362
yr_built      0.054012
sqft_lot15    0.082447
sqft_lot      0.089661
yr_renovated  0.126434
floors        0.256794
waterfront    0.266369
lat           0.307003
bedrooms      0.308797
sqft_basement 0.323816
view          0.397293
bathrooms     0.525223
sqft_living15 0.585379
sqft_above    0.605567
grade         0.667434
sqft_living   0.702035
price         1.000000
Name: price, dtype: float64
```

4 Module 4: Model Development

We can Fit a linear regression model using the longitude feature 'long' and calculate the R^2 .

```
[51]: X = df[['long']]
Y = df['price']
lm = LinearRegression()
lm.fit(X,Y)
lm.score(X, Y)
```

```
[51]: 0.00046769430149007363
```

4.0.1 Question 6

Fit a linear regression model to predict the 'price' using the feature 'sqft_living' then calculate the R^2 . Take a screenshot of your code and the value of the R^2 . You will need to submit it for the final project.

```
[52]: X= df[['sqft_living']]
      Y= df['price']
      lm= LinearRegression()
      lm.fit(X,Y)
      lm.score(X,Y)
```

```
[52]: 0.4928532179037931
```

4.0.2 Question 7

Fit a linear regression model to predict the 'price' using the list of features:

```
[64]: features = df[["floors", "waterfront","lat" ,"bedrooms" ,"sqft_basement"␣
      ↪,"view" ,"bathrooms","sqft_living15","sqft_above","grade","sqft_living"]]
```

Then calculate the R^2 . Take a screenshot of your code and the value of the R^2 . You will need to submit it for the final project.

```
[65]: lm.fit(features,Y)
      lm.score(features,Y)
```

```
[65]: 0.6576153966053727
```

4.0.3 This will help with Question 8

Create a list of tuples, the first element in the tuple contains the name of the estimator:

'scale'

'polynomial'

'model'

The second element in the tuple contains the model constructor

StandardScaler()

PolynomialFeatures(include_bias=False)

LinearRegression()

```
[67]: Input = [('scale', StandardScaler()), ('polynomial',␣
      ↪PolynomialFeatures(include_bias = False)), ('model', LinearRegression())]
```

4.0.4 Question 8

Use the list to create a pipeline object to predict the 'price', fit the object using the features in the list features, and calculate the R^2 . Take a screenshot of your code and the value of the R^2 . You will need to submit it for the final project.

```
[70]: pipe = Pipeline(Input)
      features = features.astype(float)
      pipe.fit(features, Y)
      ypipe = pipe.predict(features)
      r2_score(Y, ypipe)
```

```
[70]: 0.7499839618515916
```

5 Module 5: Model Evaluation and Refinement

```
[71]: from sklearn.model_selection import cross_val_score
      from sklearn.model_selection import train_test_split
```

We will split the data into training and testing sets:

```
[72]: features = ["floors", "waterfront","lat" ,"bedrooms" ,"sqft_basement" ,"view"
      ↪,"bathrooms","sqft_living15","sqft_above","grade","sqft_living"]
      X = df[features]
      Y = df['price']

      x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size=0.15,
      ↪random_state=1)

      print("number of test samples:", x_test.shape[0])
      print("number of training samples:",x_train.shape[0])
```

```
number of test samples: 3242
```

```
number of training samples: 18371
```

5.0.1 Question 9

Create and fit a Ridge regression object using the training data, set the regularization parameter to 0.1, and calculate the R^2 using the test data. Take a screenshot of your code and the value of the R^2 . You will need to submit it for the final project.

```
[73]: from sklearn.linear_model import Ridge

      ridge_model = Ridge(alpha=0.1)
      ridge_model.fit(x_train, y_train)
      yridge = ridge_model.predict(x_test)
      r2_score(y_test, yridge)
```

[73]: 0.6478659006870546

5.0.2 Question 10

Perform a second order polynomial transform on both the training data and testing data. Create and fit a Ridge regression object using the training data, set the regularisation parameter to 0.1, and calculate the R^2 utilising the test data provided. Take a screenshot of your code and the R^2 . You will need to submit it for the final project.

```
[76]: poly = PolynomialFeatures(degree = 2)
      x_train_poly = poly.fit_transform(x_train)
      x_test_poly = poly.fit_transform(x_test)
      ridge_model.fit(x_train_poly, y_train)
      yhat = ridge_model.predict(x_test_poly)
      r2_score(y_test, yhat)
```

[76]: 0.7001498445907909

[]: